

```
In [1]: # Importing the necessary libraries
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```

```
In [2]: # Loading the dataset
df = pd.read_csv('med dataset.csv')
```

```
In [3]: df.head()
```

```
Out[3]:
```

	id	age	year	sex	glang	part	job	stud_h	health	psyt	jspe	qcae_cog	qcae_aff	amsp	erec_mean	ci
0	2	18	1	1	120	1	0	56	3	0	88	62	27	17	0.738095	
1	4	26	4	1	1	1	0	20	4	0	109	55	37	22	0.690476	
2	9	21	3	2	1	0	0	36	3	0	106	64	39	17	0.690476	
3	10	21	2	2	1	0	1	51	5	0	101	52	33	18	0.833333	
4	13	21	3	1	1	1	0	22	4	0	102	58	28	21	0.690476	

```
In [4]: # column types
df.dtypes
```

```
Out[4]:
```

id	int64
age	int64
year	int64
sex	int64
glang	int64
part	int64
job	int64
stud_h	int64
health	int64
psyt	int64
jspe	int64
qcae_cog	int64
qcae_aff	int64
amsp	int64
erec_mean	float64
cesd	int64
stai_t	int64
mbi_ex	int64
mbi_cy	int64
mbi_ea	int64
dtype:	object

```
In [5]: # checking for missing values
df.isnull().sum() # no missing values
```

```
Out[5]:
```

id	0
age	0
year	0
sex	0
glang	0
part	0
job	0
stud_h	0
health	0
psyt	0
jspe	0
qcae_cog	0
qcae_aff	0

```
amsp      0
erec_mean  0
cesd       0
stai_t     0
mbi_ex     0
mbi_cy     0
mbi_ea     0
dtype: int64
```

```
In [6]: # removing the 'id' column
df.drop(['id'], axis=1, inplace=True)
```

```
In [7]: # check for duplicates
df.duplicated().sum() # no duplicates
```

```
Out[7]: 0
```

```
In [8]: # separate the data into two groups: categorical and numerical
df_cat = df[['sex', 'year', 'part', 'glang', 'job', 'stud_h', 'health', 'psyt']]
df_num = df[['age', 'jspe', 'qcae_aff', 'amsp', 'erec_mean', 'cesd', 'stai_t', 'mbi_ex',
```

```
In [9]: # Categorical Data
# Description
df_cat.head(20)
```

```
Out[9]:
```

	sex	year	part	glang	job	stud_h	health	psyt
0	1	1	1	120	0	56	3	0
1	1	4	1	1	0	20	4	0
2	2	3	0	1	0	36	3	0
3	2	2	0	1	1	51	5	0
4	1	3	1	1	0	22	4	0
5	2	5	1	1	1	10	2	0
6	2	5	1	1	0	15	3	0
7	1	4	1	1	1	8	4	0
8	2	4	1	1	1	20	2	0
9	2	2	1	1	0	20	5	0
10	1	1	0	1	0	20	4	0
11	2	1	0	1	1	9	5	1
12	2	5	0	90	0	25	5	0
13	2	2	1	1	0	51	2	1
14	2	2	1	1	0	42	3	0
15	1	2	1	1	0	40	4	0
16	1	1	1	90	0	9	4	0
17	1	4	1	1	0	10	4	0
18	2	5	0	1	0	38	4	0
19	2	3	1	1	0	15	4	0

```
In [10]: # Descriptive Statistics
df_cat.describe(include="all")
```

	sex	year	part	glang	job	stud_h	health	psyt
--	-----	------	------	-------	-----	--------	--------	------

Out[10]:

count	886.000000	886.000000	886.000000	886.000000	886.000000	886.000000	886.000000	886.000000
mean	1.695260	3.102709	0.563205	14.327314	0.348758	25.288939	3.777652	0.224605
std	0.472665	1.763937	0.496269	32.366389	0.476847	15.927875	1.061497	0.417558
min	1.000000	1.000000	0.000000	1.000000	0.000000	0.000000	1.000000	0.000000
25%	1.000000	1.000000	0.000000	1.000000	0.000000	12.000000	3.000000	0.000000
50%	2.000000	3.000000	1.000000	1.000000	0.000000	25.000000	4.000000	0.000000
75%	2.000000	5.000000	1.000000	1.000000	1.000000	36.000000	5.000000	0.000000
max	3.000000	6.000000	1.000000	121.000000	1.000000	70.000000	5.000000	1.000000

In [11]: *# Plotting the count of each category for each feature using Seaborn*

Set the figure size

```
plt.figure(figsize=(20, 20))
```

Plot the count of each category for each feature

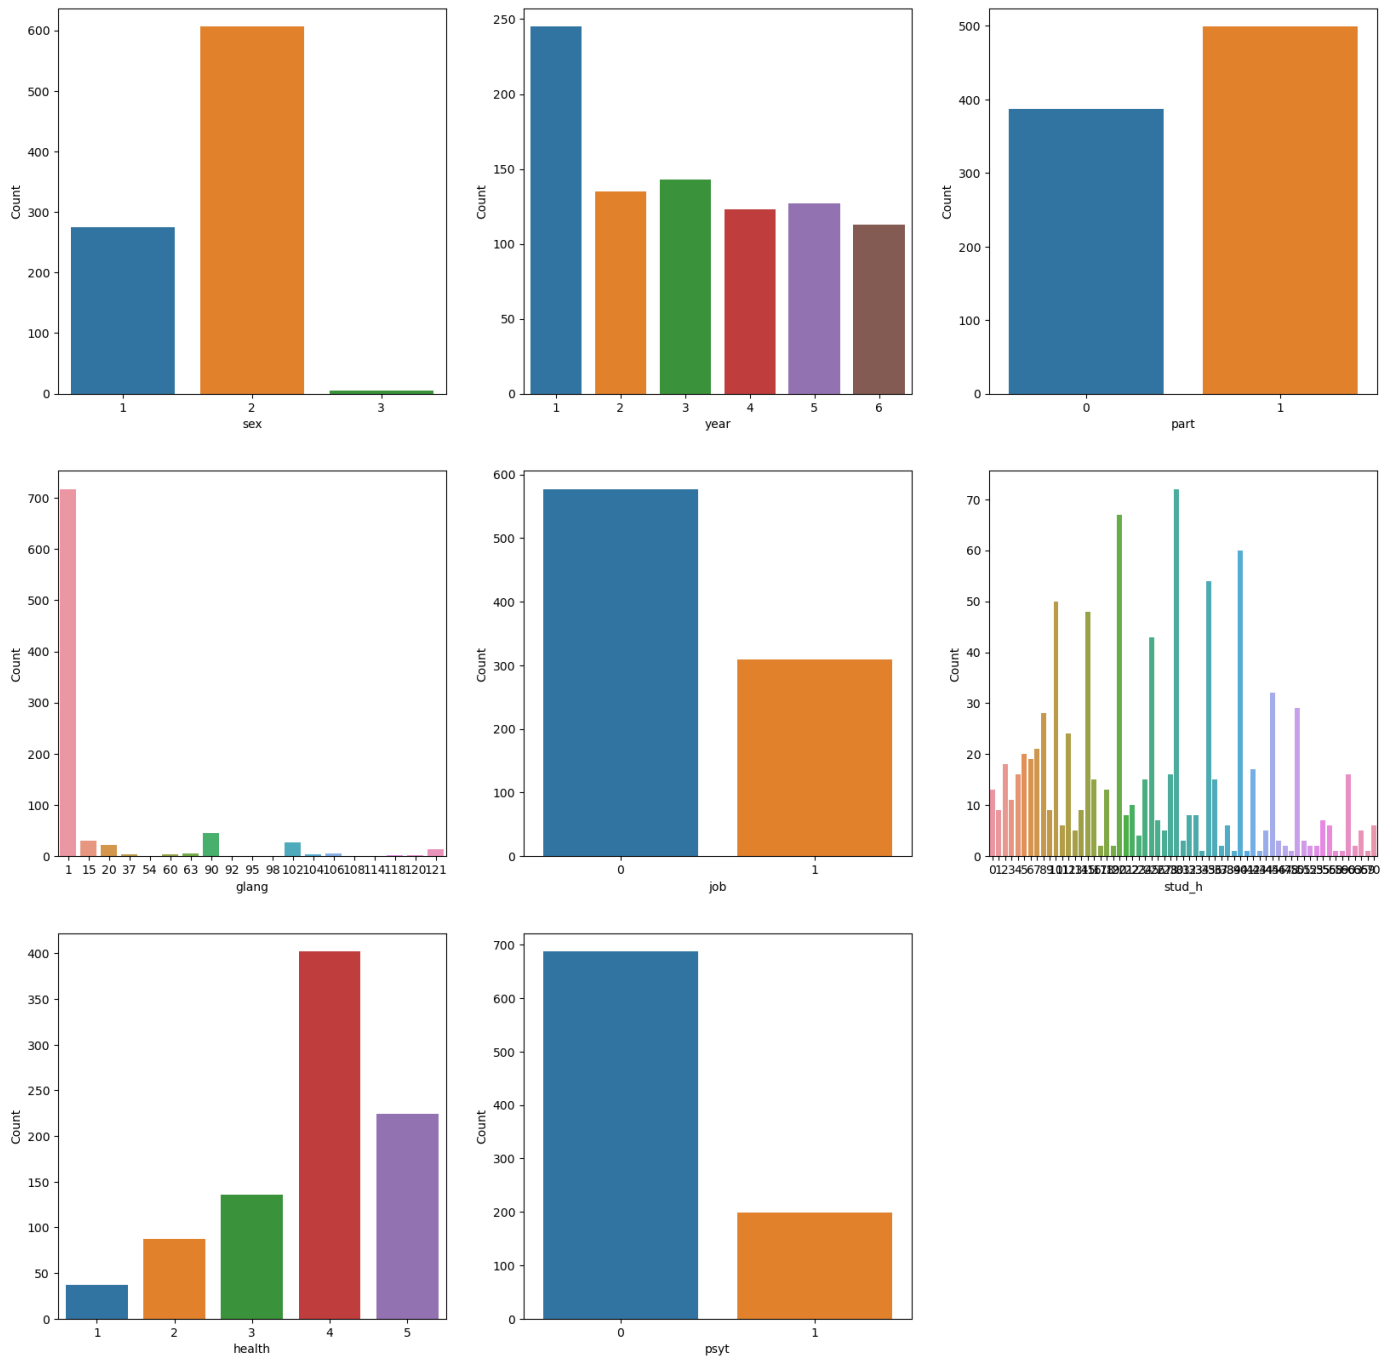
```
for i, col in enumerate(df_cat.columns):
```

```
    plt.subplot(3, 3, i+1)
```

```
    sns.countplot(x=col, data=df_cat)
```

```
    plt.xlabel(col)
```

```
    plt.ylabel('Count')
```



```
In [12]: # Finding relationships between categorical variables using Chi-Square test
from scipy.stats import chi2_contingency
# create an empty list to store the results
chi2_results = []

#Loop through each pair of categorical columns and perform Chi-Square test
for i, col1 in enumerate(df_cat.columns):
    for j, col2 in enumerate(df_cat.columns):
        if i < j:
            chi2, p, dof, ex = chi2_contingency(pd.crosstab(df_cat[col1], df_cat[col2]))
            chi2_results.append({'Variable 1': col1, 'Variable 2': col2, 'Chi-Square': c

# Create a DataFrame from the results list
chi2_table = pd.DataFrame(chi2_results)

# Display the resulting DataFrame
print(chi2_table)
```

	Variable 1	Variable 2	Chi-Square	p-value
0	sex	year	12.769894	2.368234e-01
1	sex	part	6.787509	3.358236e-02
2	sex	glang	54.444023	2.496473e-02

3	sex	job	0.399311	8.190127e-01
4	sex	stud_h	102.535577	8.736027e-01
5	sex	health	21.975553	4.961357e-03
6	sex	psyt	23.484975	7.948815e-06
7	year	part	27.375938	4.818767e-05
8	year	glang	109.823807	7.634917e-02
9	year	job	72.560233	3.002212e-14
10	year	stud_h	608.992141	3.320205e-23
11	year	health	21.604790	3.623350e-01
12	year	psyt	4.315674	5.049176e-01
13	part	glang	22.638009	2.048861e-01
14	part	job	1.811389	1.783416e-01
15	part	stud_h	65.031867	3.058703e-01
16	part	health	7.025088	1.345683e-01
17	part	psyt	0.515159	4.729141e-01
18	glang	job	19.719717	3.488184e-01
19	glang	stud_h	1113.757550	2.317115e-01
20	glang	health	84.094422	1.559493e-01
21	glang	psyt	7.207154	9.882601e-01
22	job	stud_h	81.825491	3.207564e-02
23	job	health	5.115215	2.756787e-01
24	job	psyt	2.909025	8.808513e-02
25	stud_h	health	293.764489	1.013227e-02
26	stud_h	psyt	42.095944	9.616637e-01
27	health	psyt	27.583952	1.514423e-05

```
In [13]: # Keep only the pairs of variables that are related (p-value < 0.05) and sort them by p-
chi2_table = chi2_table[chi2_table['p-value'] < 0.05].sort_values(by='p-value')
chi2_table
#from the results, we can conclude that:
# sex and health are both related to psychoterapy in the last year
# sex and health are related
# how long students study is related to their health
```

```
Out[13]:
```

	Variable 1	Variable 2	Chi-Square	p-value
10	year	stud_h	608.992141	3.320205e-23
9	year	job	72.560233	3.002212e-14
6	sex	psyt	23.484975	7.948815e-06
27	health	psyt	27.583952	1.514423e-05
7	year	part	27.375938	4.818767e-05
5	sex	health	21.975553	4.961357e-03
25	stud_h	health	293.764489	1.013227e-02
2	sex	glang	54.444023	2.496473e-02
22	job	stud_h	81.825491	3.207564e-02
1	sex	part	6.787509	3.358236e-02

```
In [14]: # Percentage of each category for each pair of variables in chi2_table
for i, row in chi2_table.iterrows():
    var1 = row['Variable 1']
    var2 = row['Variable 2']
    print(var1 + " and " + var2)
    print(pd.crosstab(index=df[var1], columns=df[var2], normalize=True))
    print('-----')
```

```
year and stud_h
stud_h      0      1      2      3      4      5      6  \
year
1      0.001129  0.000000  0.000000  0.000000  0.001129  0.001129  0.006772
```

2	0.0000000	0.0000000	0.0000000	0.0000000	0.001129	0.0000000	0.001129
3	0.002257	0.0000000	0.002257	0.001129	0.0000000	0.004515	0.001129
4	0.001129	0.001129	0.001129	0.002257	0.004515	0.005643	0.003386
5	0.002257	0.0000000	0.003386	0.001129	0.001129	0.005643	0.001129
6	0.007901	0.009029	0.013544	0.007901	0.010158	0.005643	0.007901

stud_h	7	8	9	...	53	55	56 \
year				...			
1	0.009029	0.006772	0.004515	...	0.001129	0.005643	0.005643
2	0.002257	0.001129	0.0000000	...	0.001129	0.002257	0.001129
3	0.003386	0.004515	0.0000000	...	0.0000000	0.0000000	0.0000000
4	0.006772	0.004515	0.002257	...	0.0000000	0.0000000	0.0000000
5	0.001129	0.006772	0.001129	...	0.0000000	0.0000000	0.0000000
6	0.001129	0.007901	0.002257	...	0.0000000	0.0000000	0.0000000

stud_h	58	59	60	63	65	69	70
year							
1	0.001129	0.001129	0.010158	0.002257	0.003386	0.001129	0.005643
2	0.0000000	0.0000000	0.004515	0.0000000	0.001129	0.0000000	0.001129
3	0.0000000	0.0000000	0.001129	0.0000000	0.0000000	0.0000000	0.0000000
4	0.0000000	0.0000000	0.0000000	0.0000000	0.001129	0.0000000	0.0000000
5	0.0000000	0.0000000	0.002257	0.0000000	0.0000000	0.0000000	0.0000000
6	0.0000000	0.0000000	0.0000000	0.0000000	0.0000000	0.0000000	0.0000000

[6 rows x 61 columns]

year and job

job	0	1
year		
1	0.223476	0.053047
2	0.118510	0.033860
3	0.092551	0.068849
4	0.064334	0.074492
5	0.071106	0.072235
6	0.081264	0.046275

sex and psyt

psyt	0	1
sex		
1	0.272009	0.038375
2	0.498871	0.185102
3	0.004515	0.001129

health and psyt

psyt	0	1
health		
1	0.033860	0.007901
2	0.060948	0.037246
3	0.106095	0.047404
4	0.356659	0.097065
5	0.217833	0.034989

year and part

part	0	1
year		
1	0.151242	0.125282
2	0.069977	0.082393
3	0.064334	0.097065
4	0.049661	0.089165
5	0.065463	0.077878
6	0.036117	0.091422

sex and health

health	1	2	3	4	5
sex					
1	0.014673	0.025959	0.034989	0.132054	0.102709

```
2      0.025959  0.072235  0.116253  0.320542  0.148984
3      0.001129  0.000000  0.002257  0.001129  0.001129
```

stud_h and health

```
health      1      2      3      4      5
stud_h
0      0.002257  0.003386  0.001129  0.004515  0.003386
1      0.000000  0.001129  0.000000  0.003386  0.005643
2      0.000000  0.005643  0.001129  0.009029  0.004515
3      0.000000  0.000000  0.000000  0.009029  0.003386
4      0.002257  0.001129  0.001129  0.005643  0.007901
...      ...      ...      ...      ...      ...
60     0.001129  0.002257  0.006772  0.002257  0.005643
63     0.000000  0.000000  0.001129  0.000000  0.001129
65     0.000000  0.001129  0.000000  0.002257  0.002257
69     0.000000  0.000000  0.001129  0.000000  0.000000
70     0.000000  0.002257  0.000000  0.001129  0.003386
```

[61 rows x 5 columns]

sex and glang

```
glang      1      15      20      37      54      60      63  \
sex
1      0.251693  0.010158  0.007901  0.000000  0.000000  0.000000  0.002257
2      0.554176  0.024831  0.016930  0.003386  0.001129  0.003386  0.002257
3      0.003386  0.000000  0.000000  0.000000  0.000000  0.000000  0.001129

glang      90      92      95      98      102      104      106  \
sex
1      0.022573  0.001129  0.000000  0.000000  0.005643  0.001129  0.001129
2      0.027088  0.000000  0.001129  0.001129  0.024831  0.003386  0.005643
3      0.001129  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000

glang      108      114      118      120      121
sex
1      0.001129  0.000000  0.000000  0.001129  0.004515
2      0.000000  0.001129  0.002257  0.001129  0.010158
3      0.000000  0.000000  0.000000  0.000000  0.000000
```

job and stud_h

```
stud_h      0      1      2      3      4      5      6  \
job
0      0.010158  0.004515  0.010158  0.007901  0.010158  0.009029  0.012415
1      0.004515  0.005643  0.010158  0.004515  0.007901  0.013544  0.009029

stud_h      7      8      9      ...      53      55      56  \
job
0      0.014673  0.014673  0.006772  ...  0.001129  0.007901  0.006772
1      0.009029  0.016930  0.003386  ...  0.001129  0.000000  0.000000

stud_h      58      59      60      63      65      69      70
job
0      0.001129  0.001129  0.015801  0.002257  0.004515  0.001129  0.004515
1      0.000000  0.000000  0.002257  0.000000  0.001129  0.000000  0.002257
```

[2 rows x 61 columns]

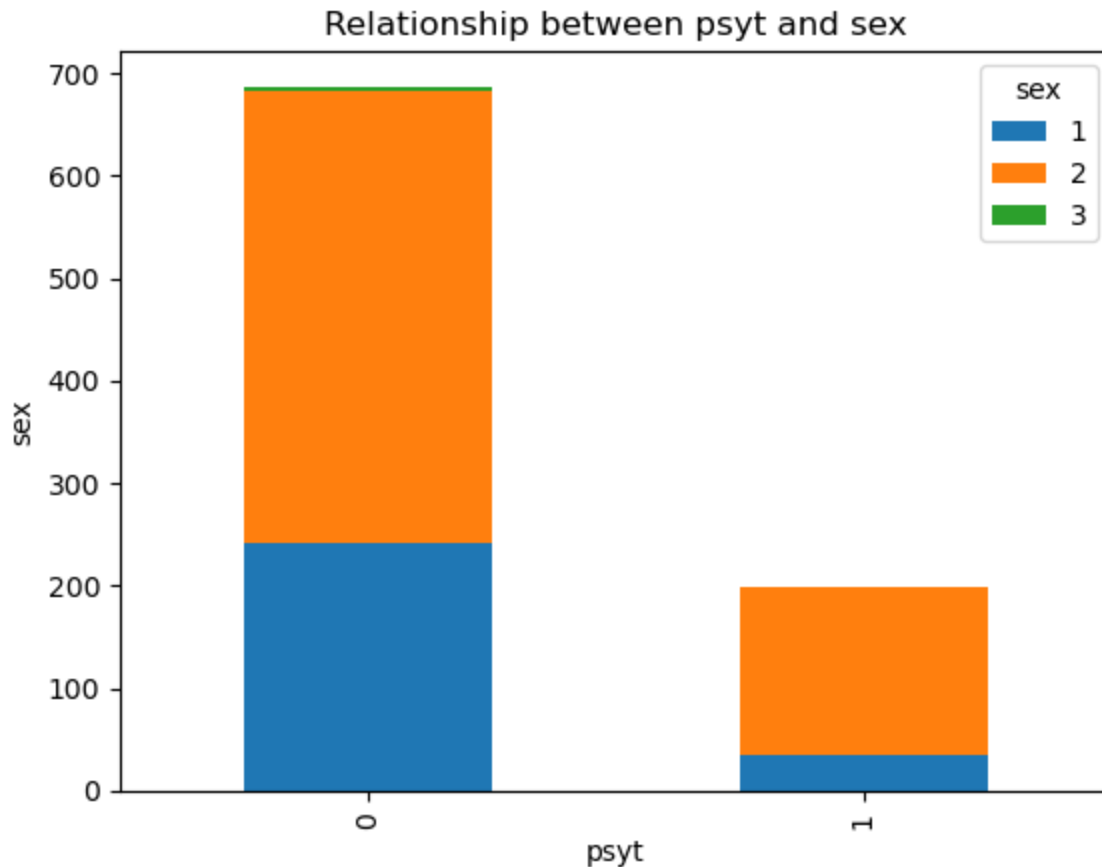
sex and part

```
part      0      1
sex
1      0.138826  0.171558
2      0.292325  0.391648
3      0.005643  0.000000
```

```
In [15]: # relationships between psyt and sex
var1 = "psyt"
var2 = "sex"

# create a cross-tabulation table
ctab = pd.crosstab(df[var1], df[var2])
```

```
In [16]: # plot the cross-tabulation table
ctab.plot(kind='bar', stacked = True)
plt.title("Relationship between " + var1 + " and " + var2)
plt.xlabel(var1)
plt.ylabel(var2)
plt.show()
```



```
In [17]: # let's see what sex visited a psychoterapist the most in the last year

# create a new data frame
df_psynt = pd.DataFrame(data = {'psyt' : df_cat['psyt'], 'sex' : df_cat['sex']})

#group the data by psychotherapy and sex
df_psynt = df_psynt.groupby(['psyt', 'sex']).size().reset_index(name='count')

# calculate the male, female and non binary in each health category
total = df_psynt['count'].sum()
df_psynt['percentage'] = df_psynt['count']/total *100
df_psynt['sex'] = df_psynt['sex'].map({1: 'Male', 2: 'Female', 3: 'Non-binary'})
```

```
In [18]: df_psynt # females tend to visit the psychoterapist more than males or non-binary people
```

```
Out[18]:
```

	psyt	sex	count	percentage
0	0	Male	241	27.200903
1	0	Female	442	49.887133
2	0	Non-binary	4	0.451467

3	1	Male	34	3.837472
4	1	Female	164	18.510158
5	1	Non-binary	1	0.112867

```
In [19]: # print each percentage for every pair of variables in chi2_table
for i, row in chi2_table.iterrows():
    var1 = row['Variable 1']
    var2 = row['Variable 2']
    print(f"{var1} and {var2}")

    # Create a temporary DataFrame with the pair of variables
    df_temp = pd.DataFrame(data={var1: df_cat[var1], var2: df_cat[var2]})
    df_temp = df_temp.groupby([var1, var2]).size().reset_index(name='count')
    total = df_temp['count'].sum()
    df_temp['percentage'] = df_temp['count'] / total * 100
    # Append the sum of numeric columns
    sum_row = df_temp.sum(numeric_only=True)
    sum_row[var1] = 'Total'
    sum_row[var2] = ''
    df_temp = pd.concat([df_temp, pd.DataFrame(sum_row).T], ignore_index=True)

    print(df_temp)
    print('-----')
```

year and stud_h

	year	stud_h	count	percentage
0	1	0	1	0.112867
1	1	4	1	0.112867
2	1	5	1	0.112867
3	1	6	6	0.677201
4	1	7	8	0.902935
..
218	6	36	2	0.225734
219	6	40	3	0.3386
220	6	45	1	0.112867
221	6	50	1	0.112867
222	Total		886.0	100.0

[223 rows x 4 columns]

year and job

	year	job	count	percentage
0	1	0	198	22.34763
1	1	1	47	5.30474
2	2	0	105	11.851016
3	2	1	30	3.386005
4	3	0	82	9.255079
5	3	1	61	6.884876
6	4	0	57	6.433409
7	4	1	66	7.44921
8	5	0	63	7.110609
9	5	1	64	7.223476
10	6	0	72	8.126411
11	6	1	41	4.62754
12	Total		886.0	100.0

sex and psyt

	sex	psyt	count	percentage
0	1	0	241	27.200903
1	1	1	34	3.837472
2	2	0	442	49.887133
3	2	1	164	18.510158
4	3	0	4	0.451467

5	3	1	1	0.112867
6	Total		886.0	100.0

health and psyt

	health	psyt	count	percentage
0	1	0	30	3.386005
1	1	1	7	0.790068
2	2	0	54	6.094808
3	2	1	33	3.724605
4	3	0	94	10.609481
5	3	1	42	4.740406
6	4	0	316	35.665914
7	4	1	86	9.706546
8	5	0	193	21.783296
9	5	1	31	3.498871
10	Total		886.0	100.0

year and part

	year	part	count	percentage
0	1	0	134	15.124153
1	1	1	111	12.528217
2	2	0	62	6.997743
3	2	1	73	8.239278
4	3	0	57	6.433409
5	3	1	86	9.706546
6	4	0	44	4.96614
7	4	1	79	8.916479
8	5	0	58	6.546275
9	5	1	69	7.78781
10	6	0	32	3.611738
11	6	1	81	9.142212
12	Total		886.0	100.0

sex and health

	sex	health	count	percentage
0	1	1	13	1.467269
1	1	2	23	2.595937
2	1	3	31	3.498871
3	1	4	117	13.205418
4	1	5	91	10.27088
5	2	1	23	2.595937
6	2	2	64	7.223476
7	2	3	103	11.625282
8	2	4	284	32.054176
9	2	5	132	14.89842
10	3	1	1	0.112867
11	3	3	2	0.225734
12	3	4	1	0.112867
13	3	5	1	0.112867
14	Total		886.0	100.0

stud_h and health

	stud_h	health	count	percentage
0	0	1	2	0.225734
1	0	2	3	0.3386
2	0	3	1	0.112867
3	0	4	4	0.451467
4	0	5	3	0.3386
..
189	69	3	1	0.112867
190	70	2	2	0.225734
191	70	4	1	0.112867
192	70	5	3	0.3386
193	Total		886.0	100.0

[194 rows x 4 columns]

```

-----
sex and glang
      sex glang  count percentage
0      1      1    223    25.1693
1      1     15      9    1.015801
2      1     20      7    0.790068
3      1     63      2    0.225734
4      1     90     20    2.257336
5      1     92      1    0.112867
6      1    102      5    0.564334
7      1    104      1    0.112867
8      1    106      1    0.112867
9      1    108      1    0.112867
10     1    120      1    0.112867
11     1    121      4    0.451467
12     2      1    491   55.417607
13     2     15     22    2.48307
14     2     20     15    1.693002
15     2     37      3    0.3386
16     2     54      1    0.112867
17     2     60      3    0.3386
18     2     63      2    0.225734
19     2     90     24    2.708804
20     2     95      1    0.112867
21     2     98      1    0.112867
22     2    102     22    2.48307
23     2    104      3    0.3386
24     2    106      5    0.564334
25     2    114      1    0.112867
26     2    118      2    0.225734
27     2    120      1    0.112867
28     2    121      9    1.015801
29     3      1      3    0.3386
30     3     63      1    0.112867
31     3     90      1    0.112867
32  Total      886.0    100.0

```

```

-----
job and stud_h
      job stud_h  count percentage
0      0      0      9    1.015801
1      0      1      4    0.451467
2      0      2      9    1.015801
3      0      3      7    0.790068
4      0      4      9    1.015801
..     ...     ...     ...     ...
101     1     53      1    0.112867
102     1     60      2    0.225734
103     1     65      1    0.112867
104     1     70      2    0.225734
105  Total      886.0    100.0

```

[106 rows x 4 columns]

```

-----
sex and part
      sex part  count percentage
0      1     0    123   13.882619
1      1     1    152   17.155756
2      2     0    259   29.232506
3      2     1    347   39.164786
4      3     0      5    0.564334
5  Total      886.0    100.0

```

```

In [20]: # Numerical variables
         df_num.head()

```

```
Out[20]:
```

	age	jspe	qcae_aff	amsp	erec_mean	cesd	stai_t	mbi_ex	mbi_cy	mbi_ea
0	18	88	27	17	0.738095	34	61	17	13	20
1	26	109	37	22	0.690476	7	33	14	11	26
2	21	106	39	17	0.690476	25	73	24	7	23
3	21	101	33	18	0.833333	17	48	16	10	21
4	21	102	28	21	0.690476	14	46	22	14	23

```
In [21]: # Distribution of the numerical variables
plt.figure(figsize = (20,20))
for i, col in enumerate(df_num.columns):
    plt.subplot(4, 3, i+1)
    sns.distplot(df_num[col])
    plt.xlabel(col)
    plt.ylabel('Density')
    plt.show()
```


C:\Users\Oana\AppData\Local\Temp\ipykernel_15732\2063736890.py:5: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df_num[col])
```




C:\Users\Oana\AppData\Local\Temp\ipykernel_15732\2063736890.py:5: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df_num[col])
```




C:\Users\Oana\AppData\Local\Temp\ipykernel_15732\2063736890.py:5: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df_num[col])
```




C:\Users\Oana\AppData\Local\Temp\ipykernel_15732\2063736890.py:5: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df_num[col])
```




C:\Users\Oana\AppData\Local\Temp\ipykernel_15732\2063736890.py:5: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df_num[col])
```




C:\Users\Oana\AppData\Local\Temp\ipykernel_15732\2063736890.py:5: UserWarning:

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df_num[col])
```




C:\Users\Oana\AppData\Local\Temp\ipykernel_15732\2063736890.py:5: UserWarning:

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df_num[col])
```




C:\Users\Oana\AppData\Local\Temp\ipykernel_15732\2063736890.py:5: UserWarning:

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df_num[col])
```




C:\Users\Oana\AppData\Local\Temp\ipykernel_15732\2063736890.py:5: UserWarning:

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df_num[col])
```




C:\Users\Oana\AppData\Local\Temp\ipykernel_15732\2063736890.py:5: UserWarning:

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df_num[col])
```




```
In [22]: # Correlation matrix
corr_matrix = df_num.corr(method = 'pearson')
corr_matrix
```

	age	jspe	qcae_aff	amsp	erec_mean	cesd	stai_t	mbi_ex	mbi_cy
age	1.000000	0.223221	-0.008130	0.000477	-0.018699	-0.139106	-0.081893	-0.182733	-0.001853
jspe	0.223221	1.000000	0.263384	0.099395	0.097062	-0.080058	-0.075036	-0.040542	-0.006981
qcae_aff	-0.008130	0.263384	1.000000	-0.071391	0.141379	0.250947	0.331350	0.215886	0.128488
amsp	0.000477	0.099395	-0.071391	1.000000	0.002780	-0.152052	-0.249231	-0.073011	-0.029343
erec_mean	-0.018699	0.097062	0.141379	0.002780	1.000000	0.029881	0.037688	0.015348	0.061965
cesd	-0.139106	-0.080058	0.250947	-0.152052	0.029881	1.000000	0.715728	0.605617	0.407727

stai_t	-0.081893	-0.075036	0.331350	-0.249231	0.037688	0.715728	1.000000	0.530486	0.331884
mbi_ex	-0.182733	-0.040542	0.215886	-0.073011	0.015348	0.605617	0.530486	1.000000	0.505200
mbi_cy	-0.001853	-0.006981	0.128488	-0.029343	0.061965	0.407727	0.331884	0.505200	1.000000
mbi_ea	0.046130	0.082508	-0.113891	0.220616	-0.034889	-0.453589	-0.462535	-0.480821	-0.565939

```
In [23]: # heatmap of the correlation matrix
plt.figure(figsize = (10,10))
sns.heatmap(corr_matrix, annot = True, cmap = 'coolwarm')
plt.show()
```



```
In [24]: # get a dataframe of the the most to least correlated variables
corr_matrix = corr_matrix.unstack().reset_index()
corr_matrix.columns = ['Variable 1', 'Variable 2', 'Correlation']
corr_matrix = corr_matrix[corr_matrix['Variable 1'] != corr_matrix['Variable 2']]
corr_matrix = corr_matrix.sort_values(by = 'Correlation', ascending = False)
corr_matrix
```


Out[24]:

	Variable 1	Variable 2	Correlation
65	stai_t	cesd	0.715728
56	cesd	stai_t	0.715728
75	mbi_ex	cesd	0.605617
57	cesd	mbi_ex	0.605617
67	stai_t	mbi_ex	0.530486
...
96	mbi_ea	stai_t	-0.462535
79	mbi_ex	mbi_ea	-0.480821
97	mbi_ea	mbi_ex	-0.480821
89	mbi_cy	mbi_ea	-0.565939
98	mbi_ea	mbi_cy	-0.565939

90 rows × 3 columns

```
In [25]: # Principal Component Analysis (PCA)
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
```

```
In [26]: # standardize the data
scaler = StandardScaler()
df_num_scaled = scaler.fit_transform(df_num)

# create a PCA instance
pca = PCA(n_components = 0.95)

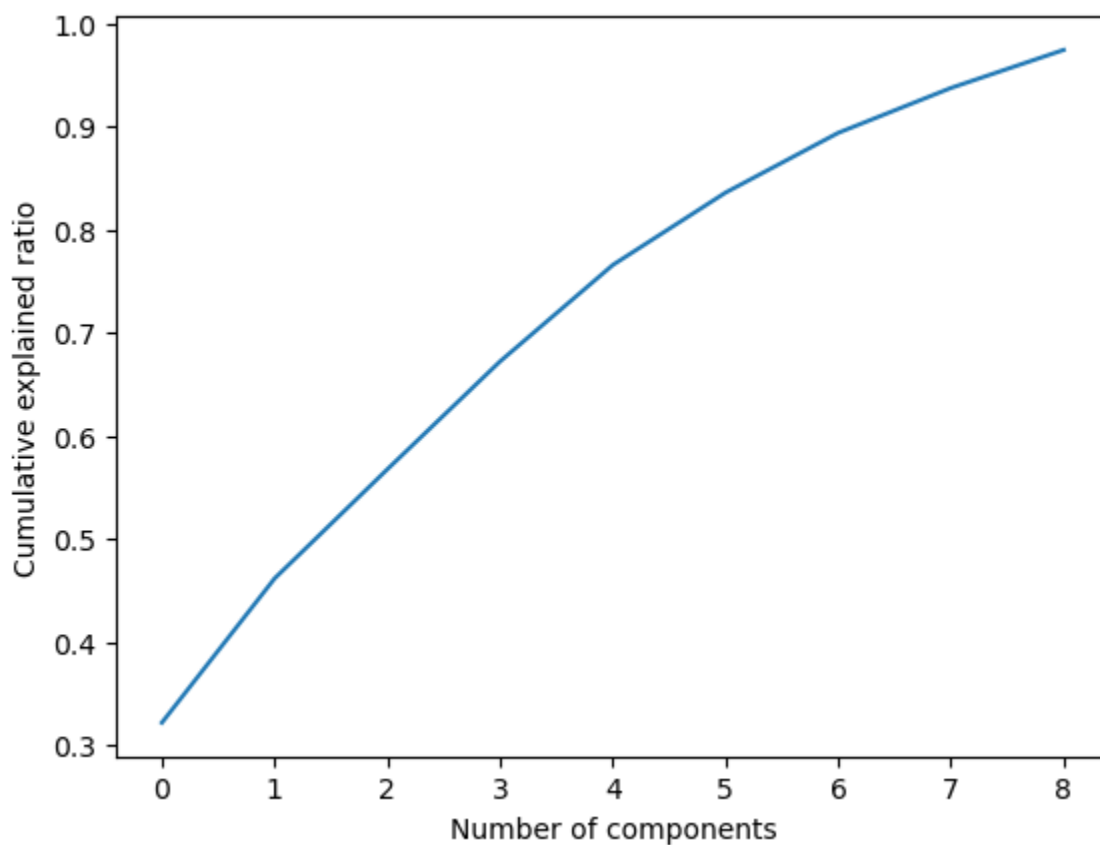
# fit the PCA instance to the scaled samples
pca.fit(df_num_scaled)

# transform the scaled samples
pca_features = pca.transform(df_num_scaled)

# print the shape of the pca features
print("The shape of the PCA is: ", pca_features.shape)
```

The shape of the PCA is: (886, 9)

```
In [27]: # plot the cumulative sum of the explained variance ratio
plt.plot(np.cumsum(pca.explained_variance_ratio_))
plt.xlabel('Number of components')
plt.ylabel('Cumulative explained ratio')
plt.show()
# the explained variance tells how much information can be attributed to each of the pri
```



```
In [28]: # loadings = coefficients that describe how each variable contributes to the principal c
loadings = pd.DataFrame({'Feature' : df_num.columns, 'PC1' : pca.components_[0], 'PC2':
loadings
```

```
Out[28]:
```

	Feature	PC1	PC2	PC3	PC4
0	age	-0.096826	0.383863	0.700289	-0.143540
1	jspe	-0.046286	0.680183	0.093058	0.065010
2	qcae_aff	0.208184	0.509245	-0.318282	-0.204569
3	amsp	-0.151479	0.121522	-0.109618	0.827293
4	erec_mean	0.045743	0.326902	-0.455469	0.052569
5	cesd	0.461591	-0.028029	-0.079636	-0.026066
6	stai_t	0.450715	0.022358	-0.074328	-0.221556
7	mbi_ex	0.442653	-0.033443	-0.030296	0.220002
8	mbi_cy	0.373686	0.036472	0.299290	0.378758
9	mbi_ea	-0.410257	0.072821	-0.278935	-0.067228

```
In [29]: # sort the loadings in descending order
loadings = loadings.sort_values(by = ['PC1', 'PC2', 'PC3', 'PC4'], ascending = False)
loadings.head(20)
# for the first pc, the variables with the highest absolute loadings are 'cesd', 'stai_t
# for the second pc, the variables with the highest absolute loadings is 'jspe'
# for the third pc - 'erec_mean'
# for the fourth pc - 'amsp'
```

```
Out[29]:
```

	Feature	PC1	PC2	PC3	PC4
5	cesd	0.461591	-0.028029	-0.079636	-0.026066
6	stai_t	0.450715	0.022358	-0.074328	-0.221556
7	mbi_ex	0.442653	-0.033443	-0.030296	0.220002

8	mbi_cy	0.373686	0.036472	0.299290	0.378758
2	qcae_aff	0.208184	0.509245	-0.318282	-0.204569
4	erec_mean	0.045743	0.326902	-0.455469	0.052569
1	jspe	-0.046286	0.680183	0.093058	0.065010
0	age	-0.096826	0.383863	0.700289	-0.143540
3	amsp	-0.151479	0.121522	-0.109618	0.827293
9	mbi_ea	-0.410257	0.072821	-0.278935	-0.067228

```
In [30]: # Clustering
from sklearn.cluster import KMeans
```


```
In [32]: # determine the optimal number of clusters
# create a list of inertia values for different k values
inertia = []
for k in range(1,10):
    # create a KMeans instance with k clusters
    model = KMeans(n_clusters = k)
    #fit model to samples
    model.fit(pca_features)
    #append the inertia to the list of inertias
    inertia.append(model.inertia_)

# plot the ks vs inertias
plt.plot(range(1,10), inertia, '-o')
plt.xlabel('Number of clusters, k')
plt.ylabel('Inertia')
plt.show()

# the point where the inertia begins to decrease more slowly is a good choice for the n
```

```
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning:
The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=4.
  warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning:
The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=4.
  warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning:
The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=4.
  warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning:
The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
  warnings.warn(
```


```
warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less chunks th
an available threads. You can avoid it by setting the environment variable OMP_NUM_THREA
DS=4.
    warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning:
The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_
init` explicitly to suppress the warning
    warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less chunks th
an available threads. You can avoid it by setting the environment variable OMP_NUM_THREA
DS=4.
    warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning:
The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_
init` explicitly to suppress the warning
    warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less chunks th
an available threads. You can avoid it by setting the environment variable OMP_NUM_THREA
DS=4.
    warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning:
The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_
init` explicitly to suppress the warning
    warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less chunks th
an available threads. You can avoid it by setting the environment variable OMP_NUM_THREA
DS=4.
    warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning:
The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_
init` explicitly to suppress the warning
    warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less chunks th
an available threads. You can avoid it by setting the environment variable OMP_NUM_THREA
DS=4.
    warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning:
The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_
init` explicitly to suppress the warning
    warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less chunks th
an available threads. You can avoid it by setting the environment variable OMP_NUM_THREA
DS=4.
    warnings.warn(
```



```
In [34]: # we choose n=3 clusters
model = KMeans(n_clusters = 3)
clusters = model.fit_predict(pca_features)
```

```
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:870: FutureWarning:
The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_
init` explicitly to suppress the warning
  warnings.warn(
C:\Users\Oana\anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1382: UserWarning:
KMeans is known to have a memory leak on Windows with MKL, when there are less chunks th
an an available threads. You can avoid it by setting the environment variable OMP_NUM_THREA
DS=4.
  warnings.warn(
```

```
In [35]: # create a scatter plot of the first two principal components
plt.scatter(pca_features[:, 0], pca_features[:, 1], c=clusters, cmap = 'viridis')
plt.xlabel('PC1')
plt.ylabel('PC2')
plt.show()
```



```
In [36]: # extracting the cluster labels
clusters_labels = model.labels_
```

```
In [37]: # adding the cluster labels to the original data frame
df_clustered = df_num.copy()
df_clustered['Cluster'] = clusters_labels
df_clustered.head()
```

```
Out[37]:
```

	age	jspe	qcae_aff	amsp	erec_mean	cesd	stai_t	mbi_ex	mbi_cy	mbi_ea	Cluster
0	18	88	27	17	0.738095	34	61	17	13	20	0
1	26	109	37	22	0.690476	7	33	14	11	26	1
2	21	106	39	17	0.690476	25	73	24	7	23	2
3	21	101	33	18	0.833333	17	48	16	10	21	0
4	21	102	28	21	0.690476	14	46	22	14	23	0

```
In [38]: # grouping the data frame by cluster to get the properties of each cluster
cluster_grouped = df_clustered.groupby('Cluster')
cluster_properties = cluster_grouped.mean()
print(cluster_properties)
```

	age	jspe	qcae_aff	amsp	erec_mean	cesd	\
Cluster							
0	20.737410	99.460432	32.798561	22.514388	0.691076	18.165468	
1	23.639769	110.198847	34.106628	24.688761	0.723823	9.481268	
2	22.467433	108.655172	37.800766	21.781609	0.746214	29.321839	


	stai_t	mbi_ex	mbi_cy	mbi_ea
Cluster				
0	43.348921	17.208633	9.708633	23.798561
1	34.354467	13.097983	7.593660	27.268012
2	53.777778	21.551724	13.777778	20.574713

```
In [39]: cluster_grouped.size()
```

```
Out[39]: Cluster
0      278
1      347
2      261
dtype: int64
```

```
In [41]: # plotting the properties for each cluster
cluster_properties.plot(kind = 'bar', figsize = (15,10))
plt.show()

# Cluster 0 is the cluster with the lowest values of the variables; the patients in this
# Cluster 1 is the cluster with the intermediate values of the variables; the patients i
# Cluster 2 is the cluster with the highest values of the variables; the patients in thi
```



```
In [42]: # Machine Learning
# Logistic Regression
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

```
In [44]: # split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(pca_features, clusters_labels, test_
```

```
In [45]: # create a logistic regression classifier
logreg = LogisticRegression()
```

```
In [46]: # fit the classifier to the training data
logreg.fit(X_train, y_train)
```

```
Out[46]: LogisticRegression
LogisticRegression()
```

```
In [47]: # predict the labels of the test set
y_pred = logreg.predict(X_test)
```

```
In [49]: # compute and print the accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy: {}".format(accuracy))
```

Accuracy: 0.9887640449438202

```
In [50]: # Model evaluation
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
```

```
In [51]: print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.97	0.98	58
1	0.98	1.00	0.99	65
2	0.98	1.00	0.99	55
accuracy			0.99	178
macro avg	0.99	0.99	0.99	178
weighted avg	0.99	0.99	0.99	178

```
[[56  1  1]
 [ 0 65  0]
 [ 0  0 55]]
```

```
In [52]: # precision is the ability of the classifier not to label a positive sample that is nega
# recall is the ability of the classifier to find all positive labels
# f1-score is the weighted average of the precision and recall
# support is the number of occurrences of each class in y_test
```

```
In [ ]: # CONCLUSION
# In this project, I used the chi-square test to determine the relationships between the
# For continuous variables, I used the Pearson correlation coefficient to determine the
# data set, and I used the K-Means clustering algorithm to cluster the patients into 3 g
# Regression model to predict the cluster labels of the patients based on the principal
# labels
```